

## Workshop – Introduction into R

# Data manipulation using tidyverse

Andreas Limacher



# Tidyverse

- The tidyverse is a collection of R packages designed for data science that share a common design philosophy and syntax.
- Some core packages:
  - ggplot2 (data visualization)
  - dplyr (data manipulation)
  - tidyr (data tidying)
  - tibble (modern data frames)
- Consistent syntax: Uses a readable code style with functions that can be chained together using the pipe operator (`%>%`)
  - Shortcut: Control+Shift+M



# Pros and cons of dplyr versus base R



## ■ Pros

- dplyr is significantly faster than base R, especially for large datasets. It can be 20-100 times faster for certain operations.
- dplyr's syntax allows for function chaining, making code cleaner and easier to read and write.
- dplyr has a set of functions focused on common data manipulation tasks, making it simpler to use

## ■ Cons

- Some operations, particularly those involving row manipulations, can be simpler in base R
- New users may find it challenging to learn dplyr's syntax and approach, especially if they're already familiar with base R
- Base R doesn't require additional package installations

# Core dplyr functions



function	description
<code>select()</code>	keep or remove columns (variables)
<code>filter()</code>	keep certain rows
<code>distinct()</code>	deduplicate rows
<code>rename()</code>	rename columns
<code>mutate()</code>	create and transform columns
<code>arrange()</code>	sort rows
<code>recode()</code>	recode levels of a factor
<code>pull()</code>	extract values from a column

# Load packages and data

- Install and load packages

```
> library(tidyverse)
> library(skimr)
```

- Load and inspect data

```
> library(NHANES)
> data(NHANES)
```

- Task: Try out the function skim. What does it do?

```
> skim(NHANES)
```



# Select variables and rows

- Select single and multiple variables (columns) and entries (rows)

```
> select(NHANES, Age) %>% pull() %>% mean()
> mean(pull(select(NHANES, Age)))

> NHANES_subset <- select(NHANES, ID, SurveyYr, Gender, Age)
> NHANES_subset <- NHANES %>% select(ID, SurveyYr, Gender, Age) # alternative
> NHANES_subset <- filter(NHANES_subset, row_number() %in% 1:200)
> NHANES_subset <- NHANES %>% select(ID, SurveyYr, Gender, Age) %>%
  filter(row_number() %in% 1:200) # combined
```

## Select – helper functions

function	description
<code>everything()</code>	all other columns not mentioned
<code>last_col()</code>	the last column
<code>contains()</code>	columns containing a character string example: <code>select(contains("time"))</code>
<code>starts_with()</code>	matches to a specified prefix example: <code>select(starts_with("date_"))</code>
<code>ends_with()</code>	matches to a specified suffix example: <code>select(ends_with("_post"))</code>
<code>num_range()</code>	a numerical range like <code>x01</code> , <code>x02</code> , <code>x03</code>
<code>any_of()</code>	matches IF column exists but returns no error if it is not found

# Select, order, and remove

## ■ Re-order variables

```
> NHANES_subset <- NHANES %>% select(SurveyYr, ID, Age, everything())
```

## ■ Select variables

```
> NHANES_subset <- NHANES %>% select(ID, SurveyYr, Gender, Age,  
  starts_with(c("BMI", "BP")), contains("Income"), last_col())
```

## ■ Remove variables

```
> NHANES_subset <- NHANES_subset %>% select(!"PregnantNow")  
> NHANES_subset <- NHANES_subset %>% select(-c("BPSys3", "BPDia3"))
```



# Deduplication

- Identify and remove duplicates

```
> NHANES_unique <- NHANES %>% distinct()  
> NHANES_unique <- NHANES %>% distinct(ID, SurveyYr)
```

- Why are the resulting unique sets not of the same size?

# Rename variables

- Rename the variable – `rename(NEW = OLD)`

```
> NHANES_new <- NHANES %>%  
  rename(Year = SurveyYr,  
         Sex = Gender)
```

# Generate or modify variables

## ■ Calculate BMI

```
> NHANES_new <- NHANES %>%  
  mutate(BMI_new = Weight / ((Height/100)^2)) %>%  
  select(ID, Weight, Height, BMI, BMI_new)
```

## ■ Calculate high income

```
> NHANES_new <- NHANES %>%  
  mutate(BMI_new = Weight / ((Height/100)^2)) %>%  
  mutate(HighIncome = if_else(HHIncomeMid > 75000, 1, 0)) %>%  
  select(ID, Age, Weight, Height, BMI, BMI_new, HHIncomeMid, HighIncome)
```

## ■ Modify high income

```
> NHANES_new <- NHANES_new %>%  
  mutate(HighIncome = if_else(HHIncomeMid > 50000, 1, 0))
```

# Convert and re-code

- Convert format

```
> NHANES_new <- NHANES_new %>%  
  mutate (ID = as.character(ID),  
          HighIncome = as.factor(HighIncome))
```

- Recode variables – recode(..., OLD = NEW)

```
> NHANES_new <- NHANES_new %>%  
  mutate(HighIncome = recode(HighIncome, "0" = "No", "1" = "Yes"))
```

- Task: Calculate an indicator variable for obesity (BMI>30)
  - How many survey participants are obese?

# Categorize – case\_when()

- Generate age in 20-year bands

```
> NHANES_new <- NHANES_new %>%  
  mutate(AgeBand = case_when(  
    Age < 20 ~ "0-19",  
    Age < 40 ~ "20-39",  
    Age < 60 ~ "40-59",  
    Age < 80 ~ "60-79",  
    Age >= 80 ~ "80+"  
  ))
```

# Filter rows

- Filter subgroups


```
> NHANES_subset <- NHANES %>%  
  filter(Gender == "female")
```

- Filter out missing values

```
> NHANES_subset <- NHANES %>%  
  filter(!is.na(Education))  
  
> NHANES_subset <- NHANES %>%  
  drop_na(Education, HHIncome)
```

# Missing values

- What are pitfalls when dealing with missing values?




# Missing values



- What are pitfalls when dealing with missing values?
  - Some functions don't work if missing values are present
  - Some functions delete observations if missing values are present. If many variables are used, this can lead to big data loss.
  - Deleting missing data can introduce bias.
  - Many statistical and machine learning models cannot handle missing values.
  - Caveat: Sometimes missing values are coded as 0 or 999 and must be set to NA (or a missing category).
  - Check NAs carefully when filtering data (some functions include NAs, some exclude them such as `filter()`).
  - Check NAs carefully when merging data



# Missing values in R



- Missing values in R are typically represented by NA (Not Available)
- Use `is.na()` function to identify NA values
- `sum(is.na())` can count total NA values in a dataset
- Many R functions have an `na.rm` parameter. Setting `na.rm = TRUE` ignores NA values during calculations
  - Example: `mean(x, na.rm = TRUE)` calculates the mean excluding NA values
- How many missing values are there in the variable BMI?

# Sort data

## ■ Ascending

```
> NHANES_new <- NHANES_new %>%  
  arrange(Weight)
```

## ■ Descending

```
> NHANES_new <- NHANES_new %>%  
  arrange(desc(Weight))
```

## ■ Nested

```
> NHANES_new <- NHANES_new %>%  
  arrange(desc(AgeBand), Weight)
```

# Summary tables



- The `{gtsummary}` package provides an elegant and flexible way to create publication-ready summary tables.
- Perfect for presenting descriptive statistics, comparing group demographics (e.g., creating a Table 1 for medical journals), and more.
- Automatically detects continuous, categorical, and dichotomous variables, calculates appropriate descriptive statistics, and also includes amount of missingness in each variable.
- Customizes tables using a growing list of formatting/styling functions. Bold labels, italicize levels, and add p-value.

# Summary tables



## ■ Summary table

```
> library(gtsummary)
> table1 <- NHANES_new %>%
  tbl_summary(include = c(Age, AgeBand, BMI, HighIncome))
```

## ■ by HighIncome

```
> table1 <- NHANES_new %>%
  tbl_summary(include = c(Age, AgeBand, BMI), by = HighIncome)
  %>% add_p()
```

## ■ save as HTML

```
> library(gt)
> gtsave(as_gt(table1), filename = "Table1.html")
```

# Exercise



Characteristic	Overall N = 3,722 <sup>1</sup>	female N = 1,882 <sup>1</sup>	male N = 1,840 <sup>1</sup>	p-value <sup>2</sup>
Age	40 (30, 52)	40 (30, 51)	41 (30, 52)	0.7
Race				0.015
Black	581 (16%)	317 (17%)	264 (14%)	
Hispanic	284 (7.6%)	151 (8.0%)	133 (7.2%)	
Mexican	448 (12%)	199 (11%)	249 (14%)	
White	2,066 (56%)	1,033 (55%)	1,033 (56%)	
Other	343 (9.2%)	182 (9.7%)	161 (8.8%)	
CivilStatus				0.10
Single	1,435 (39%)	750 (40%)	685 (37%)	
Partner	2,287 (61%)	1,132 (60%)	1,155 (63%)	
Education				<0.001
8th Grade	245 (6.6%)	113 (6.0%)	132 (7.2%)	
9 - 11th Grade	481 (13%)	215 (11%)	266 (14%)	
High School	779 (21%)	371 (20%)	408 (22%)	
Some College	1,171 (31%)	614 (33%)	557 (30%)	
College Grad	1,046 (28%)	569 (30%)	477 (26%)	